

Logistic_regression

March 8, 2020

```
[52]: from IPython.display import Image
      from sklearn.externals.six import StringIO
      from sklearn.tree import export_graphviz
      import pydotplus
      import imblearn
      import lightgbm
      import hyperopt
```

```
[30]: import pandas as pd
      import numpy as np
      import matplotlib
      import matplotlib.pyplot as plt
      import seaborn as sns
      import statsmodels.api as sm
      %matplotlib inline
      import plusmodules as pm
      import warnings
      warnings.filterwarnings('ignore')
```

```
[31]: df=pd.read_csv('telco_chrun_encoded (1).csv')
      df.head()
```

```
[31]:
```

| | gender | senior | partner | dependents | tenure | phone_service | multiple_lines | \ |
|---|--------|--------|---------|------------|--------|---------------|----------------|---|
| 0 | 1 | 0 | 1 | 0 | 1 | 0 | 0 | |
| 1 | 0 | 0 | 0 | 0 | 34 | 1 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 2 | 1 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 45 | 0 | 0 | |
| 4 | 1 | 0 | 0 | 0 | 2 | 1 | 0 | |

| | online_security | online_backup | device_protection | ... | total_charges | \ |
|---|-----------------|---------------|-------------------|-----|---------------|---|
| 0 | 0 | 1 | 0 | ... | 29.85 | |
| 1 | 1 | 0 | 1 | ... | 1889.50 | |
| 2 | 1 | 1 | 0 | ... | 108.15 | |
| 3 | 1 | 0 | 1 | ... | 1840.75 | |
| 4 | 0 | 0 | 0 | ... | 151.65 | |

| | churn | avg_monthly_charges | internet_service-fiber_optic | \ |
|--|-------|---------------------|------------------------------|---|
|--|-------|---------------------|------------------------------|---|

| | | | |
|---|---|-----------|---|
| 0 | 0 | 29.850000 | 0 |
| 1 | 0 | 55.573529 | 0 |
| 2 | 1 | 54.075000 | 0 |
| 3 | 0 | 40.905556 | 0 |
| 4 | 1 | 75.825000 | 1 |

| | internet_service-no | contract-one_year | contract-two_year | \ |
|---|---------------------|-------------------|-------------------|---|
| 0 | 0 | 0 | 0 | |
| 1 | 0 | 1 | 0 | |
| 2 | 0 | 0 | 0 | |
| 3 | 0 | 1 | 0 | |
| 4 | 0 | 0 | 0 | |

| | payment_method-credit_card_auto | payment_method-electronic_check | \ |
|---|---------------------------------|---------------------------------|---|
| 0 | 0 | 1 | |
| 1 | 0 | 0 | |
| 2 | 0 | 0 | |
| 3 | 0 | 0 | |
| 4 | 0 | 1 | |

| | payment_method-mailed_check |
|---|-----------------------------|
| 0 | 0 |
| 1 | 1 |
| 2 | 1 |
| 3 | 0 |
| 4 | 0 |

[5 rows x 25 columns]

```
[32]: df.shape
```

```
[32]: (7032, 25)
```

```
[33]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                7032 non-null  int64
1   senior                7032 non-null  int64
2   partner               7032 non-null  int64
3   dependents            7032 non-null  int64
4   tenure                7032 non-null  int64
5   phone_service         7032 non-null  int64
6   multiple_lines        7032 non-null  int64
```

```

7  online_security          7032 non-null  int64
8  online_backup            7032 non-null  int64
9  device_protection        7032 non-null  int64
10 tech_support             7032 non-null  int64
11 streaming_tv             7032 non-null  int64
12 streaming_movies         7032 non-null  int64
13 paperless_billing        7032 non-null  int64
14 monthly_charges          7032 non-null  float64
15 total_charges            7032 non-null  float64
16 churn                    7032 non-null  int64
17 avg_monthly_charges      7032 non-null  float64
18 internet_service-fiber_optic 7032 non-null  int64
19 internet_service-no      7032 non-null  int64
20 contract-one_year        7032 non-null  int64
21 contract-two_year        7032 non-null  int64
22 payment_method-credit_card_auto 7032 non-null  int64
23 payment_method-electronic_check 7032 non-null  int64
24 payment_method-mailed_check 7032 non-null  int64
dtypes: float64(3), int64(22)
memory usage: 1.3 MB

```

```
[6]: df=df.sample(frac=0.25, random_state=3)
```


```
[7]: y=df['churn']
     x=df.drop('churn', axis=1)
```

```
[8]: df['churn'].value_counts()
```

```
[8]: 0    1318
     1     440
     Name: churn, dtype: int64
```

```
[9]: df['churn'].value_counts(normalize=True).plot.bar()
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x29d66b51c18>
```



#Building the predictive algorithm using Random forest

```
[10]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30,
↪ random_state=3)

[11]: from sklearn.preprocessing import StandardScaler

ss= StandardScaler()
xs=ss.fit_transform(x)

x_trains =ss.fit_transform(x_train)
x_tests =ss.transform(x_test)

[12]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, roc_auc_score, accuracy_score,
↪ roc_curve

rfc=RandomForestClassifier(n_estimators=100)

rfc.fit(x_trains, y_train)

y_train_pred=rfc.predict(x_trains)

y_train_prob=rfc.predict_proba(x_trains)[: ,1]
```

```

print('Confusion Matrix - Train: ', '\n' ,confusion_matrix(y_train,
→y_train_pred))

print('Overall Accuracy - Train: ', accuracy_score(y_train, y_train_pred))

print('AUC- Train:' , roc_auc_score(y_train, y_train_prob))

y_test_pred= rfc.predict(x_tests)

y_test_prob=rfc.predict_proba(x_tests)[: ,1]

print('Confusion Matrix - Test: ', '\n' ,confusion_matrix(y_test, y_test_pred))

print('Overall Accuracy - Test: ', accuracy_score(y_test, y_test_pred))

print('AUC- Test:' , roc_auc_score(y_test, y_test_prob))

fpr, tpr, thresholds =roc_curve(y_test, y_test_prob)
plt.plot(fpr, tpr)
plt.plot(fpr, fpr, 'red')
plt.xlabel('FPR')
plt.ylabel('TPR')


```

```

Confusion Matrix - Train:
[[912  0]
 [ 0 318]]
Overall Accuracy - Train:  1.0
AUC- Train: 1.0
Confusion Matrix - Test:
[[364  42]
 [ 70  52]]
Overall Accuracy - Test:  0.7878787878787878
AUC- Test: 0.8298372769118953

```

```
[12]: Text(0, 0.5, 'TPR')
```



#Hyper parameter tuning for Random forest

```
[13]: from sklearn.model_selection import RandomizedSearchCV

from scipy.stats import randint as sp_randint

rfc=RandomForestClassifier(random_state=3)

params={'n_estimators': sp_randint(50, 200),
        'max_features': sp_randint(1, 24),
        'max_depth': sp_randint(2,10),
        'min_samples_split': sp_randint(2, 20),
        'min_samples_leaf': sp_randint(1, 20),
        'criterion': ['gini', 'entropy'] }

rsearch_rfc = RandomizedSearchCV(rfc, param_distributions= params, n_iter=50,
    ↪cv=3, scoring='roc_auc',
        random_state=3, return_train_score=True)

rsearch_rfc.fit(xs, y)

[13]: RandomizedSearchCV(cv=3, error_score=nan,
        estimator=RandomForestClassifier(bootstrap=True,
        ccp_alpha=0.0,
```

```

class_weight=None,
criterion='gini',
max_depth=None,
max_features='auto',
max_leaf_nodes=None,
max_samples=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
min_samples_split=2,

min_weight_fraction_leaf=0.0,

n_estimators=100,
n_jobs...
'min_samples_leaf':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x0000029D674FFDD8>,
'min_samples_split':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x0000029D674FFA90>,
'n_estimators':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x0000029D674FF5C0>},
pre_dispatch='2*n_jobs', random_state=3, refit=True,
return_train_score=True, scoring='roc_auc', verbose=0)

```

```
[14]: rsearch_rfc.best_params_
```

```
[14]: {'criterion': 'entropy',
      'max_depth': 8,
      'max_features': 2,
      'min_samples_leaf': 16,
      'min_samples_split': 5,
      'n_estimators': 120}
```

```
[15]: pd.DataFrame(rsearch_rfc.cv_results_).head(2)
```

```
[15]:
```

| | mean_fit_time | std_fit_time | mean_score_time | std_score_time | \ |
|---|---------------|--------------|-----------------|----------------|---|
| 0 | 0.175137 | 0.023511 | 0.012330 | 0.001515 | |
| 1 | 0.815924 | 0.052228 | 0.046503 | 0.007779 | |

| | param_criterion | param_max_depth | param_max_features | param_min_samples_leaf | \ |
|---|-----------------|-----------------|--------------------|------------------------|---|
| 0 | gini | 2 | 4 | 9 | |
| 1 | entropy | 5 | 11 | 12 | |

| | param_min_samples_split | param_n_estimators | ... | split1_test_score | \ |
|---|-------------------------|--------------------|-----|-------------------|---|
| 0 | 2 | 71 | ... | 0.830451 | |
| 1 | 11 | 188 | ... | 0.820061 | |

| | split2_test_score | mean_test_score | std_test_score | rank_test_score | \ |
|---|-------------------|-----------------|----------------|-----------------|---|
| 0 | 0.848930 | 0.841708 | 0.008065 | 10 | |

| | | | | |
|---|----------|----------|----------|----|
| 1 | 0.855198 | 0.841571 | 0.015390 | 11 |
|---|----------|----------|----------|----|

| | split0_train_score | split1_train_score | split2_train_score | \ |
|---|--------------------|--------------------|--------------------|---|
| 0 | 0.852283 | 0.860699 | 0.848861 | |
| 1 | 0.888735 | 0.896357 | 0.883738 | |

| | mean_train_score | std_train_score |
|---|------------------|-----------------|
| 0 | 0.853948 | 0.004974 |
| 1 | 0.889610 | 0.005189 |

[2 rows x 22 columns]

```
[16]: # model using best params
rfc=RandomForestClassifier(**rsearch_rfc.best_params_, random_state=3)

rfc.fit(x_trains, y_train)

y_train_pred=rfc.predict(x_trains)

y_train_prob=rfc.predict_proba(x_trains)[: ,1]

print('Confusion Matrix - Train: ', '\n' ,confusion_matrix(y_train,
↪y_train_pred))

print('Overall Accuracy - Train: ', accuracy_score(y_train, y_train_pred))

print('AUC- Train:' , roc_auc_score(y_train, y_train_prob))


y_test_pred= rfc.predict(x_tests)

y_test_prob=rfc.predict_proba(x_tests)[: ,1]

print('Confusion Matrix - Test: ', '\n' ,confusion_matrix(y_test, y_test_pred))

print('Overall Accuracy - Test: ', accuracy_score(y_test, y_test_pred))

print('AUC- Test:' , roc_auc_score(y_test, y_test_prob))


fpr, tpr, thresholds =roc_curve(y_test, y_test_prob)
plt.plot(fpr, tpr)
plt.plot(fpr, fpr, 'red')
plt.xlabel('FPR')
plt.ylabel('TPR')
```


Confusion Matrix - Train:


```

[[887 25]
 [193 125]]
Overall Accuracy - Train: 0.8227642276422764
AUC- Train: 0.8874510371841554
Confusion Matrix - Test:
[[383 23]
 [ 86 36]]
Overall Accuracy - Test: 0.7935606060606061
AUC- Test: 0.8407494145199063

```

```
[16]: Text(0, 0.5, 'TPR')
```



```

[17]: imp = pd.DataFrame(rfc.feature_importances_, index = x.columns, columns=['imp'])

imp = imp.sort_values(by = 'imp', ascending = False)

imp

# the results says the percentage of dropping based on the features, we can use
↳ this as feature selection if we remove
#the features which are having the value of 0 we can simply remove those this
↳ says that the features are not having any
#impact on data

```

```
[17]:                                     imp
      tenure                             0.155242
      total_charges                       0.119885
      contract-two_year                   0.092182
      monthly_charges                     0.085751
      avg_monthly_charges                 0.083130
      internet_service-fiber_optic       0.077960
      payment_method-electronic_check    0.076173
      online_security                     0.040575
      internet_service-no                 0.040186
      contract-one_year                   0.035663
      dependents                          0.031578
      paperless_billing                   0.023652
      online_backup                       0.023328
      payment_method-credit_card_auto    0.019833
      tech_support                        0.019323
      partner                             0.014272
      streaming_movies                    0.010863
      gender                              0.008537
      device_protection                   0.008298
      streaming_tv                        0.008101
      senior                              0.007751
      multiple_lines                      0.007676
      payment_method-mailed_check        0.007300
      phone_service                       0.002743
```

```
[18]: from sklearn.neighbors import KNeighborsClassifier

      from sklearn.model_selection import RandomizedSearchCV

      from scipy.stats import randint as sp_randint

      knn = KNeighborsClassifier()

      params= {'n_neighbors': sp_randint(1,20),
               'p': sp_randint(1,5)}

      rsearch_knn=RandomizedSearchCV(knn, param_distributions = params, cv=3,
      ↪random_state=3, n_iter=50)

      rsearch_knn.fit(xs, y)
```

```
[18]: RandomizedSearchCV(cv=3, error_score=nan,
                        estimator=KNeighborsClassifier(algorithm='auto',
                                                         leaf_size=30,
                                                         metric='minkowski',
                                                         metric_params=None,
```

```

n_jobs=None, n_neighbors=5,
p=2, weights='uniform'),
iid='deprecated', n_iter=50, n_jobs=None,
param_distributions={'n_neighbors':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x0000029D685FE1D0>,
'p':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x0000029D685FE198>},
pre_dispatch='2*n_jobs', random_state=3, refit=True,
return_train_score=False, scoring=None, verbose=0)

```

```
[19]: rsearch_knn.best_params_
```

```
[19]: {'n_neighbors': 15, 'p': 1}
```

```

[20]: knn = KNeighborsClassifier(** rsearch_knn.best_params_)
from sklearn.metrics import confusion_matrix, roc_auc_score, accuracy_score,
    ↪ roc_curve, classification_report

knn.fit(x_trains, y_train)

y_train_pred=knn.predict(x_trains)

y_train_prob=knn.predict_proba(x_trains)[:,:1]

print('Confusion Matrix - Train: ', '\n' ,confusion_matrix(y_train,
    ↪ y_train_pred))

print('Overall Accuracy - Train: ', accuracy_score(y_train, y_train_pred))

print('Classification Report - Test', classification_report(y_test,
    ↪ y_test_pred))

print('AUC- Train:' , roc_auc_score(y_train, y_train_prob))


y_test_pred= knn.predict(x_tests)

y_test_prob=knn.predict_proba(x_tests)[:,:1]

print('Confusion Matrix - Test: ', '\n' ,confusion_matrix(y_test, y_test_pred))

print('Overall Accuracy - Test: ', accuracy_score(y_test, y_test_pred))

print('AUC- Test:' , roc_auc_score(y_test, y_test_prob))

fpr, tpr, thresholds =roc_curve(y_test, y_test_prob)

```

```
plt.plot(fpr, tpr)
plt.plot(fpr, fpr, 'red')
plt.xlabel('FPR')
plt.ylabel('TPR')
```

Confusion Matrix - Train:

```
[[804 108]
```

```
[113 205]]
```

Overall Accuracy - Train: 0.8203252032520325

Classification Report - Test

| | | precision | recall | f1-score | support |
|--|--|-----------|--------|----------|---------|
|--|--|-----------|--------|----------|---------|

| | | | | |
|---|------|------|------|-----|
| 0 | 0.82 | 0.94 | 0.88 | 406 |
|---|------|------|------|-----|

| | | | | |
|---|------|------|------|-----|
| 1 | 0.61 | 0.30 | 0.40 | 122 |
|---|------|------|------|-----|

| | | | | |
|----------|--|--|------|-----|
| accuracy | | | 0.79 | 528 |
|----------|--|--|------|-----|

| | | | | |
|-----------|------|------|------|-----|
| macro avg | 0.71 | 0.62 | 0.64 | 528 |
|-----------|------|------|------|-----|

| | | | | |
|--------------|------|------|------|-----|
| weighted avg | 0.77 | 0.79 | 0.77 | 528 |
|--------------|------|------|------|-----|

AUC- Train: 0.8661487642061128

Confusion Matrix - Test:


```
[[348 58]
```

```
[ 53 69]]
```

Overall Accuracy - Test: 0.7897727272727273

AUC- Test: 0.8143018654607121

[20]: Text(0, 0.5, 'TPR')



1 stacking results of 3 learners (random forest, knn, logistic regression)

```
[21]: from sklearn.linear_model import LogisticRegression
lr= LogisticRegression(solver= 'liblinear')

lr.fit(x_trains, y_train)

y_train_pred=lr.predict(x_trains)

y_train_prob=lr.predict_proba(x_trains)[: ,1]

print('Confusion Matrix - Train: ', '\n' ,confusion_matrix(y_train,
    ↪y_train_pred))

print('Overall Accuracy - Train: ', accuracy_score(y_train, y_train_pred))

print('Classification Report - Test', classification_report(y_test,
    ↪y_test_pred))

print('AUC- Train:' , roc_auc_score(y_train, y_train_prob))

y_test_pred= lr.predict(x_tests)

y_test_prob=.predict_proba(x_tests)[: ,1]

print('Confusion Matrix - Test: ', '\n' ,confusion_matrix(y_test, y_test_pred))

print('Overall Accuracy - Test: ', accuracy_score(y_test, y_test_pred))

print('AUC- Test:' , roc_auc_score(y_test, y_test_prob))

fpr, tpr, thresholds =roc_curve(y_test, y_test_prob)
plt.plot(fpr, tpr)
plt.plot(fpr, fpr, 'red')
plt.xlabel('FPR')
plt.ylabel('TPR')
```

Confusion Matrix - Train:
[[826 86]


```
[146 172]]
Overall Accuracy - Train: 0.8113821138211382
Classification Report - Test
precision    recall  f1-score   support

      0       0.87       0.86       0.86       406
      1       0.54       0.57       0.55       122

 accuracy          0.79          528
 macro avg         0.71          528
weighted avg         0.79          528

AUC- Train: 0.8558320920225091
Confusion Matrix - Test:
[[359  47]
 [ 57  65]]
Overall Accuracy - Test: 0.803030303030303
AUC- Test: 0.843111523863361
```

```
[21]: Text(0, 0.5, 'TPR')
```



```
[22]: from sklearn.ensemble import VotingClassifier

lr = LogisticRegression(solver='liblinear')
knn = KNeighborsClassifier(**rsearch_knn.best_params_)
```

```
rfc = RandomForestClassifier(**rsearch_rfc.best_params_)
```

2 HARD VOTING

```
[24]: clf = VotingClassifier(estimators=[('lr', lr), ('knn', knn), ('rfc', rfc)],  
    ↪voting='hard')  
  
clf.fit(x_trains, y_train)  
  
y_train_pred = clf.predict(x_trains)  
y_test_pred = clf.predict(x_tests)  
  
print('Accuracy score - Train', accuracy_score(y_train, y_train_pred))  
print('Accuracy score - Test', accuracy_score(y_test, y_test_pred))
```

Accuracy score - Train 0.8195121951219512

Accuracy score - Test 0.8011363636363636

3 SOFT VOTING - Equal weightages

```
[25]: clf = VotingClassifier(estimators=[('lr', lr), ('knn', knn), ('rfc', rfc)],  
    ↪voting='soft')  
  
clf.fit(x_trains, y_train)  
  
y_train_pred=clf.predict(x_trains)  
  
y_train_prob=clf.predict_proba(x_trains)[: ,1]  
  
print('Confusion Matrix - Train: ', '\n' ,confusion_matrix(y_train,  
    ↪y_train_pred))  
  
print('Overall Accuracy - Train: ', accuracy_score(y_train, y_train_pred))  
  
print('Classification Report - Test', classification_report(y_test,  
    ↪y_test_pred))  
  
print('AUC- Train:' , roc_auc_score(y_train, y_train_prob))  
  
y_test_pred= clf.predict(x_tests)  
  
y_test_prob=clf.predict_proba(x_tests)[: ,1]
```

```

print('Confusion Matrix - Test: ', '\n' ,confusion_matrix(y_test, y_test_pred))

print('Overall Accuracy - Test: ', accuracy_score(y_test, y_test_pred))

print('AUC- Test:' , roc_auc_score(y_test, y_test_prob))

fpr, tpr, thresholds =roc_curve(y_test, y_test_prob)
plt.plot(fpr, tpr)
plt.plot(fpr, fpr, 'red')
plt.xlabel('FPR')
plt.ylabel('TPR')

```

Confusion Matrix - Train:

```
[[837  75]
```

```
[147 171]]
```

Overall Accuracy - Train: 0.8195121951219512

Classification Report - Test

| | | precision | recall | f1-score | support |
|--------------|------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.90 | 0.87 | | 406 |
| 1 | 0.59 | 0.47 | 0.52 | | 122 |
| accuracy | | | 0.80 | | 528 |
| macro avg | 0.72 | 0.68 | 0.70 | | 528 |
| weighted avg | 0.79 | 0.80 | 0.79 | | 528 |

AUC- Train: 0.8763516495641621

Confusion Matrix - Test:


```
[[365  41]
```

```
[ 66  56]]
```

Overall Accuracy - Test: 0.7973484848484849

AUC- Test: 0.8382459823952193

[25]: Text(0, 0.5, 'TPR')



```
[26]: clf = VotingClassifier(estimators=[('lr', lr), ('knn', knn), ('rfc', rfc)],
    ↳ weights=[1,2,3] ,voting='soft')

clf.fit(x_trains, y_train)

y_train_pred=clf.predict(x_trains)

y_train_prob=clf.predict_proba(x_trains)[: ,1]

print('Confusion Matrix - Train: ', '\n' ,confusion_matrix(y_train,
    ↳ y_train_pred))

print('Overall Accuracy - Train: ', accuracy_score(y_train, y_train_pred))

print('Classification Report - Test', classification_report(y_test,
    ↳ y_test_pred))

print('AUC- Train:' , roc_auc_score(y_train, y_train_prob))

y_test_pred= clf.predict(x_tests)

y_test_prob=clf.predict_proba(x_tests)[: ,1]
```

```

print('Confusion Matrix - Test: ', '\n' ,confusion_matrix(y_test, y_test_pred))

print('Overall Accuracy - Test: ', accuracy_score(y_test, y_test_pred))

print('AUC- Test:' , roc_auc_score(y_test, y_test_prob))

fpr, tpr, thresholds =roc_curve(y_test, y_test_prob)
plt.plot(fpr, tpr)
plt.plot(fpr, fpr, 'red')
plt.xlabel('FPR')
plt.ylabel('TPR')

```

Confusion Matrix - Train:

```

[[849  63]
 [153 165]]

```

Overall Accuracy - Train: 0.824390243902439

Classification Report - Test precision recall f1-score
support

| | | | | |
|---|------|------|------|-----|
| 0 | 0.85 | 0.90 | 0.87 | 406 |
| 1 | 0.58 | 0.46 | 0.51 | 122 |

| | | | | |
|--------------|------|------|------|-----|
| accuracy | | | 0.80 | 528 |
| macro avg | 0.71 | 0.68 | 0.69 | 528 |
| weighted avg | 0.78 | 0.80 | 0.79 | 528 |

AUC- Train: 0.8806307238221339

Confusion Matrix - Test:

```

[[365  41]
 [ 69  53]]

```

Overall Accuracy - Test: 0.7916666666666666

AUC- Test: 0.8369538883953807

[26]: Text(0, 0.5, 'TPR')

